

brinkgrad: An open-source FEniCSx framework for adjoint-based topology optimisation of coupled flow-transport in porous media using Brinkman–convection-diffusion equations

Nisong Monyimba^{1*}, Vincent Pizziconi^{1†}, Aurel Coza^{1‡}

¹School of Biological and Health Systems Engineering,

Arizona State University, Tempe, AZ, USA

`nisongmonyimba278@gmail.com`

2025

Abstract

brinkgrad is an MIT-licensed open-source FEniCSx package implementing adjoint-based topology optimisation for coupled Brinkman–convection-diffusion systems in porous media. All constituent methods are standard; the contribution is their integration into a minimal, installable package with full verification, CI, and one-command reproducibility. The package provides SUPG-stabilised forward and adjoint solvers, Helmholtz filter, Heaviside projection, and an OC optimiser within a single reproducible default preset, implemented in 1,903 lines of documented Python across 9 modules. Post-processing binary-geometry analysis is supported via `robust_projection()` (see `examples/`).

Gradient correctness is confirmed by adjoint direction and Taylor remainder tests (Section 4). The framework is validated by 18 automated tests (10 unit, 8 integration) executed on every commit via GitHub Actions inside a pinned `dolfinx/dolfinx:v0.7.3` Docker container. Seven executable example scripts demonstrate the framework across multiple target profiles; all results are reproducible from a single command. To the authors’ knowledge, no openly documented and continuously tested FEniCSx implementation combining coupled Brinkman–convection-diffusion topology optimisation, SUPG-stabilised forward and adjoint transport, CI-based verification, Docker reproducibility, and packaged distribution has been reported. The framework is released under the MIT licence with Docker-based reproducibility, continuous integration, and permanent archival through Zenodo (DOI: 10.5281/zenodo.20538833, [1]).

*ORCID: 0009-0000-7558-8580

†ORCID: to be added by co-author

‡ORCID: to be added by co-author

1 Introduction

Density-based topology optimisation for Stokes flow, introduced by Borrvall and Petersson [2], has been extended to conjugate heat transfer [3], scalar transport problems [4], and unsteady flows [5]. The Brinkman–SIMP approach replaces solid regions with a high-permeability porous medium, enabling gradient-based optimisation of fluid topology without remeshing. Microfluidic systems provide a natural benchmark application because the small length scales of soft-lithography devices make manual re-design expensive and simulation-driven optimisation increasingly accessible [6].

Despite growing interest, several computational gaps remain in the open-source literature on coupled flow-transport topology optimisation.

Gaps addressed by this work

Gap 1: Open-source implementation of coupled flow-transport topology optimisation. Existing open-source topology optimisation codes for fluidic systems target pure flow [2] or global scalar mixing [7]. FEniCSx-based frameworks exist for structural problems [8] and Stokes flow [9], but a documented, tested, and reproducible open-source implementation for the coupled Brinkman–convection-diffusion system with a boundary concentration-mismatch objective has not been published. The linear gradient benchmark is chosen for verifiability as a sanity check; the framework’s primary contribution is the implementation methodology, not the benchmark RMSE.

Gap 2: High-Péclet transport in FEniCSx TO. Existing FEniCSx topology optimisation implementations [8] operate at $Pe \ll 1$ and do not require transport stabilisation. At microfluidic Péclet numbers ($Pe \sim 10^2\text{--}10^5$), SUPG stabilisation of both forward and adjoint equations is required [10,11]; we document this standard technique in the FEniCSx UFL context.

Gap 3: End-to-end reproducibility. Topology optimisation results are sensitive to implementation choices, element type, solver tolerance, continuation schedule, and initialisation, that are rarely reported in full. Published microfluidic topology optimisation results [7,12] are not accompanied by open-source code, making independent verification and extension difficult. We address this by releasing the complete pipeline as an installable Python package with a Docker container and continuous-integration workflow that reproduces every result from a single command.

Table 1 compares documented features of **brinkgrad** against existing open-source frameworks.

Contributions

All techniques used (continuous adjoint, SUPG stabilisation, Brinkman–SIMP, Helmholtz filtering, Heaviside projection, OC, MMA) are established methods from the literature. The contribution of this paper is their integration into a complete, documented, tested, and reproducible open-source FEniCSx framework for coupled flow-transport topology optimisation. Specifically:

1. **Open-source implementation of coupled Brinkman–convection-diffusion topology optimisation in FEniCSx.** A complete, installable Python package implementing the coupled forward problem, continuous adjoint (including the coupling

Table 1: Feature comparison based on publicly documented capabilities of open-source topology optimisation frameworks. ✓ = implemented and documented; — = not reported in cited publications or public repositories. Entries reflect capabilities at the time of writing; unlisted features may exist but are not verified by the authors.

Feature		BP03 [2]	Na22 [7]	Fe24 [12]	d-adj [13]	cashocs [14]	ours
<i>Physics</i>							
Coupled	flow-transport	—	✓	✓	—	—	✓
Conc.-mismatch	objective	—	✓	—	—	—	✓
SUPG adjoint	stabilisation	—	—	—	—	—	✓
Continuous	adjoint	✓	—	✓	—	✓	✓
Discrete	adjoint	—	—	—	✓	✓	future
<i>Software engineering</i>							
FEniCSx	backend	—	—	—	✓	✓	✓
Docker	image (pinned)	—	—	—	—	—	✓
CI / tests	automated	—	—	—	—	—	✓
Zenodo	archival (DOI)	—	—	—	—	—	✓
Jupyter	examples	—	—	—	—	✓	✓
Open (MIT)	source	partial	—	—	✓	✓	✓

term $-\lambda \nabla c$ and misfit-derived Dirichlet outlet condition), and sensitivity computation. Existing frameworks (dolfin-adjoint [13], cashocs [14]) provide general adjoint capabilities but lack: (i) a pre-configured implementation of the coupled Brinkman–CD system with SUPG stabilisation at $Pe \sim 10^2$ – 10^5 ; (ii) a default configuration with failure-mode documentation; (iii) end-to-end Docker/CI reproducibility; and (iv) a concentration-mismatch objective with explicit adjoint derivation (Table 1, Section 2.7).

- 2. Documented implementation of stabilised forward and adjoint transport at $Pe \sim 10^2$ – 10^5 .** The standard SUPG stabilisation [10, 11] is applied consistently to both forward and adjoint transport equations and verified by finite-difference Pearson direction correlation ≥ 0.80 confirming correct descent direction (Section 4).
- 3. Reproducible software architecture.** Docker container, GitHub Actions CI on every commit, Zenodo archival, pytest suite, and Jupyter notebook. All results in this paper are reproducible from a single command in approximately one hour on a standard laptop (Section 3.9).
- 4. Verification and software-validation suite.** Gradient verification (Taylor remainder, direction test), mesh sensitivity study (four meshes), binary gap quantification (Supplementary S5), and OC/MMA schedule verification (pure-MMA output RMSE = 0.304 vs. hybrid RMSE = 0.079, confirming the default configuration functions as intended; Table 4).

Nature of contribution. This paper is a *software implementation paper*. The contribution is not a new optimisation algorithm, not a record-breaking RMSE. The contribution is a documented, tested, and reproducible open-source implementation of adjoint-based

topology optimisation for a coupled Brinkman–convection–diffusion system in FEniCSx; something that does not currently exist in the open-source literature.

Scope and limitations. `brinkgrad`¹ operates within the Brinkman–SIMP porous-medium surrogate on a fixed 80×20 mesh. Two open directions are identified and quantified: (i) binary-convergent design generation via robust projection (`robust_projection()`), implemented and demonstrated in `examples/robust_projection_demo.py` [15] (thresholding results: Supplementary S5); and (ii) mesh-independent continuation (as expected for non-convex topology optimisation, finer meshes may find different local optima; RMSE 0.091–0.107 on 160×40 vs. 0.058 on the 80×20 primary mesh).

The remainder of this paper is organised as follows. Section 2 presents the mathematical model and adjoint derivation. Section 3 describes the FEniCSx implementation and configuration and continuation settings (Section 3.4). Section 4 presents verification studies. Section 5 presents the benchmark application. Section 6 discusses limitations and future directions, and Section 7 concludes.

2 Mathematical Model

2.1 Design domain and boundary conditions

We consider a two-dimensional rectangular domain $\Omega = (0, L_x) \times (0, L_y)$ with $L_x = 2000 \mu\text{m}$ and $L_y = 500 \mu\text{m}$. The left boundary is split into two inlets of equal height:

- **Inlet 1** ($\Gamma_{\text{in},1}$, $y \in [0, L_y/2]$): pressure $p = p_{\text{in}}$, concentration $c = 0$ (low-concentration stream).
- **Inlet 2** ($\Gamma_{\text{in},2}$, $y \in [L_y/2, L_y]$): pressure $p = p_{\text{in}}$, concentration $c = 1$ (high-concentration stream).

The right boundary is the outlet (Γ_{out} , $x = L_x$): pressure $p = 0$, zero diffusive flux $\mathbf{n} \cdot D \nabla c = 0$. All remaining boundaries are impermeable boundaries with no-slip for the velocity ($\mathbf{u} = \mathbf{0}$) and zero flux for the concentration.

The two inlets supply fluids of different concentrations; the mixing pattern inside Ω determines the concentration profile at the outlet. The design is represented by a continuous material density field $\rho(\mathbf{x}) \in [0, 1]$, where $\rho = 1$ denotes a fully permeable (fluid) region and $\rho = 0$ a fully impermeable (solid-like) region. Intermediate values $\rho \in (0, 1)$ represent porous material with intermediate permeability and diffusivity, consistent with the Brinkman–SIMP modelling framework [2].

2.2 Brinkman–Darcy flow

Notation for the design variable. The design variable $\rho(\mathbf{x}) \in [0, 1]$ is a continuous *permeability interpolation parameter*, not a physical porosity. $\rho = 1$ corresponds to $\alpha \approx 0$ (fluid-like permeability); $\rho = 0$ corresponds to $\alpha = \alpha_{\text{max}}$ (solid-like impermeability). This parameterisation follows the SIMP-style penalisation of Borrvall & Petersson [2].

¹The package is named `brinkgrad` to denote its physical basis (Brinkman equations) and avoid namespace collisions with existing libraries (formerly `micrograd`) to avoid indexing ambiguity.

The steady, incompressible flow of a Newtonian fluid through the design domain is modelled by the Brinkman equations [2]:

$$\nabla p - \mu \nabla^2 \mathbf{u} + \alpha(\rho) \mathbf{u} = \mathbf{0}, \quad (1a)$$

$$\nabla \cdot \mathbf{u} = 0. \quad (1b)$$

Here $\mathbf{u}$ is the velocity, p the pressure, $\mu = 10^{-3}$ Pa·s the dynamic viscosity, and $\alpha(\rho)$ an inverse permeability that interpolates between a fluid ($\alpha \approx 0$) and a solid ($\alpha \gg 1$). Following Borrvall & Petersson [2], we use the convex interpolation

$$\alpha(\rho) = \alpha_{\min} + (\alpha_{\max} - \alpha_{\min}) \frac{1 - \rho}{1 + \rho}, \quad (2)$$

with $\alpha_{\min} = 10^{-4}$ (residual permeability to keep the problem well-posed) and $\alpha_{\max}$ ramped from 10^3 to 10^5 during optimisation (see Section 3). In practice, $\alpha_{\max}$ is ramped from 10^3 to 10^5 during the optimisation to avoid early stagnation in an all-solid local minimum.

Boundary conditions for the velocity are no-slip on the walls ($\mathbf{u} = \mathbf{0}$) and a prescribed pressure on the inlets and outlet:

$$p = p_{\text{in}} \text{ on } \Gamma_{\text{in},1} \cup \Gamma_{\text{in},2}, \quad p = 0 \text{ on } \Gamma_{\text{out}}. \quad (3)$$

2.3 Convection–diffusion transport

The steady-state concentration c of the solute is governed by

$$\mathbf{u} \cdot \nabla c - \nabla \cdot (D(\rho) \nabla c) = 0, \quad (4)$$

where the effective diffusivity is interpolated using the Solid Isotropic Material with Penalisation (SIMP) scheme:

$$D(\rho) = D_{\min} + (D_{\text{fluid}} - D_{\min}) \rho^p, \quad (5)$$

with $D_{\text{fluid}} = 10^{-9}$ m²/s the molecular diffusivity of the solute, $D_{\min} = 10^{-15}$ m²/s a small but non-zero lower bound (to prevent a singular diffusion matrix in solid regions; six orders of magnitude below D_{fluid} to ensure solid regions are effectively impermeable to concentration transport), and $p = 3$ the SIMP exponent.

Porous leakage and the role of $D_{\min}$. A critical question for the validity of the Brinkman–SIMP model is whether the optimiser exploits *artificial transport pathways* through gray-zone solid regions. In solid regions ($\rho \approx 0$), the inverse permeability $\alpha \rightarrow 10^5$ Pa·m⁻² (capped at 10^5 in the present implementation) suppresses flow to negligible levels ($Da \approx 4 \times 10^{-6} \ll 1$). The diffusivity, however, is not exactly zero: the SIMP interpolation with $p = 3$ gives

$$D(\rho) = D_{\min} + (D_{\text{fluid}} - D_{\min}) \rho^3,$$

so at $\rho = 0.1$, $D \approx 10^{-14}$ m² s⁻¹, six orders of magnitude below D_{fluid} . The characteristic diffusion time through one element in a solid region is $(\Delta y)^2 / D(\rho) \approx (25 \mu\text{m})^2 / 10^{-14} \approx 6 \times 10^4$ s, compared to the advective flow time $L_x / U_c \approx 20$ s. The ratio $t_{\text{diff}} / t_{\text{flow}} \approx 3000$ confirms that diffusive leakage through solid regions is negligible in steady state, even for intermediate gray densities. The present release targets Brinkman–SIMP density

optimisation and therefore inherits gray-region behaviour characteristic of density methods; quantitative analysis is in Supplementary S5. Robust projection (`robust_projection()`) is implemented and integration into the optimisation loop is in the project roadmap.

Dirichlet conditions fix the concentration at the inlets,

$$c = 1 \text{ on } \Gamma_{\text{in},2}, \quad c = 0 \text{ on } \Gamma_{\text{in},1}, \quad (6)$$

while a zero diffusive flux condition is applied at the outlet,

$$\mathbf{n} \cdot D \nabla c = 0 \text{ on } \Gamma_{\text{out}}. \quad (7)$$

No-flux conditions are imposed on the remaining walls.

2.4 Nondimensional operating regime

The operating regime is characterised by three dimensionless numbers (Table 2): Stokes flow ($Re \ll 1$), convection-dominated transport ($Pe_h \approx 780$), and valid impermeability surrogate ($Da \ll 1$).

Table 2: Operating regime for the benchmark microfluidic configuration. All values computed on the primary 80×20 mesh at peak velocity.

Parameter	Symbol	Value	Physical regime
Reynolds number	Re	3×10^{-3}	Stokes flow ($Re \ll 1$)
Cell Péclet	Pe_h	≈ 780	Convection-dominated
Darcy number	Da	4×10^{-6}	Impermeability valid ($Da \ll 1$)

The diffusion sublayer thickness at the outlet is

$$\delta \approx \sqrt{\frac{D_{\text{fluid}} L_y}{u_{\text{max}}}} \approx 4 \mu\text{m}, \quad (8)$$

which is underresolved by a factor of $\approx 6 \times$ at $\Delta y = 25 \mu\text{m}$; SUPG eliminates spurious oscillations but does not recover sub-grid diffusion structure (Section 6.6).

2.5 Optimisation problem

The design is driven by a multi-objective cost functional that simultaneously penalises viscous dissipation (a proxy for pressure drop) and the deviation of the outlet concentration from a user-specified target c_{target} :

$$J_f = \int_{\Omega} \left(\frac{\mu}{2} \|\nabla \mathbf{u}\|^2 + \frac{1}{2} \alpha(\rho) \|\mathbf{u}\|^2 \right) d\Omega, \quad (9a)$$

$$J_c = \frac{1}{2} \int_{\Gamma_{\text{out}}} (c - c_{\text{target}})^2 ds, \quad (9b)$$

$$J = w_f J_f + w_c J_c. \quad (9c)$$

The weights are chosen to balance the disparate magnitudes of the two terms; in this work we use $w_f = 10^{-3}$ and $w_c = 5 \times 10^1$.

The optimisation problem reads

$$\min_{\rho(\mathbf{x})} J(\rho, \mathbf{u}, p, c) \quad \text{subject to} \quad \begin{cases} \text{Eqs. (1) and (4)} & \text{(PDE constraints),} \\ \frac{1}{|\Omega|} \int_{\Omega} \rho \, d\Omega \leq V^* & \text{(volume constraint),} \\ 0 \leq \rho(\mathbf{x}) \leq 1 & \text{(box constraints),} \end{cases} \quad (10)$$

where $V^* = 0.5$ is the target fluid volume fraction. Note that the projected fluid volume fraction $|\{x : \bar{\rho}(x) < 0.5\}|/|\Omega| = 0.343$ differs from $1 - V^* = 0.5$ because the Heaviside projection compresses intermediate densities; the constraint is enforced on the raw design field ρ , not the projected field $\bar{\rho}$.

2.6 Design regularisation

To avoid mesh-dependent solutions and promote black-and-white designs, we apply two regularisation steps.

Helmholtz PDE filter. The design variable ρ is first smoothed by solving

$$-r_f^2 \nabla^2 \tilde{\rho} + \tilde{\rho} = \rho \quad \text{in } \Omega, \quad (11)$$

with homogeneous Neumann boundary conditions. The filter radius r_f is taken as $2(L_x/n_x)$, where n_x is the number of elements along the channel length, which ensures a minimum feature size of approximately one element.

Smoothed Heaviside projection. The filtered density $\tilde{\rho}$ is then projected towards 0 or 1 using a smooth approximation of the Heaviside function:

$$\bar{\rho} = \frac{\tanh(\beta\eta) + \tanh(\beta(\tilde{\rho} - \eta))}{\tanh(\beta\eta) + \tanh(\beta(1 - \eta))}, \quad \eta = 0.5. \quad (12)$$

The sharpness parameter β is increased during the optimisation (from 1 to 128) to gradually drive the design from a gray intermediate state to a near-binary-permeability distribution. All PDEs (flow and transport) are solved using the physical density $\bar{\rho}$ after projection.

2.7 Adjoint sensitivity analysis

We employ the continuous adjoint method [2, 9], which evaluates the total derivative $dJ/d\rho$ by solving two adjoint PDEs per optimisation iteration, independently of the number of design variables. The key result is the sensitivity expression:

$$\frac{dJ}{d\rho} = \int_{\Omega} \left[\frac{\partial \alpha}{\partial \rho} \mathbf{u} \cdot \mathbf{v} + \frac{\partial D}{\partial \rho} \nabla c \cdot \nabla \lambda \right] d\Omega, \quad (13)$$

where $(\mathbf{v}, q, \lambda)$ are the adjoint velocity, pressure, and concentration fields obtained by solving the adjoint Brinkman and adjoint convection-diffusion equations. The adjoint concentration satisfies the Dirichlet outlet condition $\lambda = c_{\text{target}} - c$ on Γ_{out} , derived directly from the concentration-mismatch objective. The full sensitivity chain through the Helmholtz filter and Heaviside projection is implemented in `adjoint.py`. Full derivation is in Supplementary SA.

3 Numerical Implementation

3.1 Finite element discretisation

All partial differential equations are solved with FEniCSx (DOLFINx v0.7.3) [16] on structured triangular meshes. The primary mesh uses 80×20 elements (1600 elements, $\Delta x = \Delta y = 25 \mu\text{m}$); a coarser 20×5 mesh is used for rapid parameter screening.

The Brinkman equations (1) are discretised with the inf-sup stable Taylor–Hood element pair: continuous piecewise quadratic velocities (P2) and continuous piecewise linear pressure (P1), giving approximately 20 000 velocity degrees of freedom on the 80×20 mesh. The convection–diffusion equation (4) and its adjoint (20) are discretised with continuous piecewise linear elements (P1) and stabilised using the SUPG method [11]. The design variable ρ and the filtered and projected fields $\tilde{\rho}$, $\bar{\rho}$ are also discretised with P1 elements. Weak forms are assembled using the Unified Form Language (UFL); all integrals are evaluated with exact quadrature.

3.2 Linear solvers

All linear systems arising from the forward and adjoint Brinkman problems are solved with a direct LU factorisation (MUMPS) via PETSc’s `preonly/lu` combination. This choice provides robust, parameter-free solutions for the moderate meshes used here ($\lesssim 30\,000$ degrees of freedom). For larger-scale problems, an iterative solver is available: a block-triangular field-split preconditioner with algebraic multigrid (HYPRE) for the velocity block and the pressure Schur complement [17]; scalability results are in Supplementary Information S1.

3.3 Optimisation algorithm

The design is updated using the Optimality Criteria (OC) method [18] as the default:

$$\rho^{\text{new}} = \max(\rho_{\min}, \min(\rho_{\max}, \rho^{\text{old}} \cdot B^\eta)), \quad B = -\frac{\text{d}J/\text{d}\rho}{\Lambda}, \quad (14)$$

with damping exponent $\eta = 0.5$ and move limit 0.15. An optional hybrid schedule (OC for iterations 0–299, MMA [19] for iterations 300–599) is also implemented and documented in Table 4, but is not the default configuration. The volume constraint is enforced explicitly in both update rules via Lagrange multiplier bisection (OC) and dual bisection (MMA).

3.4 Default configuration and continuation settings

Each component of the continuation schedule is documented in Table 3 with its motivation. The hybrid OC/MMA schedule is validated by direct comparison: pure MMA throughout stagnates near the initial objective (RMSE = 0.304, Table 4), because MMA cannot escape the uniform-density plateau without the larger OC move at $\beta = 1$. Pure OC (600 iter) achieves RMSE = 0.0053[†] on the coarse 20×5 mesh ([†] only 6 outlet nodes; under-resolves the outlet profile, so this RMSE is not directly comparable to the 80×20 primary result). Qualitatively, OC alone is sufficient for this problem. Pure OC is the default configuration; the hybrid OC/MMA schedule is retained as an optional comparator for generalisation studies. The framework is fully modular: any component can be disabled by setting its parameter to the trivial value (e.g. $r_{\text{filter}} \rightarrow 0$, $\beta \equiv 1$, pure OC or pure MMA).

Table 3: Default configuration settings: addressed failure mode, remedy, and effect on primary result. Each component is modular and independently disableable.

Setting	Failure mode addressed	Effect
Helmholtz PDE filter	Mesh-dependent checkerboard	Smooth, mesh-independent ρ_f
Heaviside projection (β -continuation)	Premature gray-zone lock-in	Drives $\bar{\rho} \rightarrow \{0, 1\}$ gradually
$\alpha_{\max}$ ramping	All-solid early stagnation	Avoids non-fluid local minima
OC optimiser	Pure OC is the default; pure MMA stagnates (Table 4); optional hybrid OC/MMA documented	Validated by three-way comparison
Sinusoidal initialisation	Bias toward uniform $\rho = V^*$	Breaks symmetry, aids branching topology

The five default settings work together to ensure reliable convergence to connected, near-thresholded-permeability designs.

Volume fraction continuation. The target fluid volume fraction V^* is fixed at 0.5 throughout all 600 iterations. Experiments with a relaxed initial $V^* = 0.7$ ramping to 0.5 over the first half of iterations were found to cause the OC update to chase a moving target, preventing convergence; fixing V^* from the start eliminates this instability.

Sinusoidal initialisation. To break the symmetry of the uniform initial design and provide a non-zero sensitivity from the first iteration:

$$\rho(\mathbf{x})|_{t=0} = 0.7 + 0.25 \sin\left(\frac{8\pi y}{L_y}\right), \quad \text{clipped to } [0.01, 0.99]. \quad (15)$$

Smooth concentration penalty. Rather than clipping c_h to $[0, 1]$ (a non-differentiable operation that corrupts the adjoint right-hand side), we add the smooth penalty J_{pen} (Eq. (27)) to the objective. This C^∞ , vanishes for $c \in [0, 1]$, and contributes a consistent term to the sensitivity.

Heaviside sharpening. The projection parameter β (Eq. (12)) is increased through the sequence

$$\beta \in \{1, 2, 4, 8, 16\}$$

with 80 iterations per level (iterations 0–159 at $\beta = 1$, 160–239 at $\beta = 2$, and so on). The gray-zone fraction at $\beta = 128$ (primary result) is 0.074; an extended 1400-iteration run with $\beta \in \{32, 64, 96, 128\}$ further reduces the gray-zone fraction to 0.074 (see Section 5.2) (Section 5.2).

Brinkman penalty ramping. $\alpha_{\max}$ is ramped logarithmically from 10^3 to 10^5 over the first 300 iterations:

$$\alpha_{\max}(t) = 10^{3+t}, \quad t = \min\left(1, 2 \cdot \frac{\text{step}}{\text{max_iter}}\right), \quad (16)$$

ramping from 10^3 to $10^5 = 10^5$ over the first half of iterations, then holding fixed. Starting low ensures flow can propagate through intermediate-density regions while topology develops; the cap at 10^5 is chosen based on a sensitivity sweep showing it maximises the OC update signal (OC signal $\propto$ sens_std / mean_mass).

Table 4: Optimiser comparison: 600-iteration run on the 20×5 mesh, linear gradient target. Pure OC is the default; hybrid OC/MMA is optional. Pure MMA stagnates; both OC-based schedules converge. All results are software verification outputs.

Schedule	RMSE	Method	Converges?	Notes
Pure MMA (600 iter)	0.304	nlopt MMA	No	stagnates at $J \approx J_0$
Pure OC (600 iter)	0.0053 [†]	OC bisect	Yes	20×5 only; see [†]
Hybrid OC/MMA (600 iter)	0.079	OC+MMA	Yes	similar to pure OC

[†] The 20×5 mesh (6 outlet nodes) under-resolves the outlet; RMSE not comparable to 80×20 results.

3.5 Definition and reporting of outlet RMSE

The outlet RMSE is defined as

$$\text{RMSE} = \sqrt{\frac{1}{N_{\text{out}}} \sum_{i=1}^{N_{\text{out}}} (c_h(y_i) - c_{\text{target}}(y_i))^2}, \quad (17)$$

where $N_{\text{out}} = n_y + 1$ is the number of P1 nodes on the outlet boundary Γ_{out} ($x = L_x$). Both c_h and c_{target} are dimensionless and normalised to $[0, 1]$. For the primary 80×20 mesh, $N_{\text{out}} = 21$ ($\Delta y = 25 \mu\text{m}$); no interpolation is applied. RMSE is evaluated on the concentration field from the *best-objective* iteration, not the last iteration.

3.6 Dimensional note

All simulations are two-dimensional. The flow rate Q and hydraulic resistance $R = \Delta p / Q$ are reported per unit channel depth (units: $\text{m}^2 \text{s}^{-1}$ and Pa m^{-2} , respectively). Three-dimensional values follow by multiplying Q by the channel depth $L_z = 100 \mu\text{m}$ (as used in the Reynolds-number calculation in Section 2.4) and dividing R by L_z .

3.7 Full optimisation loop

Algorithm 1 (Table 5) summarises the complete optimisation procedure.

Table 5: Algorithm 1: Pure OC optimisation loop with β -continuation (default configuration). Steps 2–8 repeat for $k = 0, \dots, N - 1$ where $N = 1400$ for the primary result (configurable; a 600-iteration run uses $N = 600$). An optional MMA update is available; see Section 3.3.

Step	Operation
1	Initialise $\rho \leftarrow V^*$ (sinusoidal perturbation; Eq. (15))
for $k = 0, 1, \dots, N - 1$	$(N = 1400 \text{ primary; configurable})$:
2	Look up schedule: $(\beta_k, \text{method}_k, \alpha_{\max,k}) \leftarrow \text{schedule}(k)$
3	Helmholtz filter: $\rho_f \leftarrow \text{filter}(\rho)$ [Eq. (11)]
4	Heaviside projection: $\bar{\rho} \leftarrow \text{project}(\rho_f, \beta_k)$ [Eq. (12)]
5	Forward solve: $(\mathbf{u}, p, c) \leftarrow \text{solve}(\bar{\rho}, \alpha_{\max,k})$ [Eqs. (1),(4)]
6	Adjoint solve: $(\mathbf{v}, q, \lambda) \leftarrow \text{adjoint}(\bar{\rho}, \mathbf{u}, c)$ [Section 2.7]
7	Sensitivity: $s \leftarrow dJ/d\bar{\rho}$
8	Update: $\rho \leftarrow \text{OC}(\rho, s, V^*)$ [Eq. (14); MMA optional]
9	Track best J ; store ρ_{best}
return	ρ_{best}

3.8 Practical usage and Jupyter notebook

brinkgrad is designed for rapid inverse design exploration. The package follows semantic versioning (v1.0.0); the single public entry point `GradientGeneratorOptimizer` provides a stable API with keyword arguments and documented defaults. A typical workflow proceeds in three steps:

1. **Define target:** specify the desired outlet concentration profile $c^*(y)$ as a Python lambda; no adjoint modifications are required.
2. **Run optimisation:** call `GradientGeneratorOptimizer` with mesh resolution, volume fraction V^* , and objective weights w_f, w_c ; the continuation schedule runs automatically.
3. **Post-process:** extract RMSE, gray fraction, and permeability field; optionally threshold-project to a binary design for downstream analysis.

The repository includes a Jupyter notebook (`notebooks/quickstart.ipynb`) demonstrating this workflow for the linear gradient target in under 30 lines of user code. Key API calls are:

```
from brinkgrad import GradientGeneratorOptimizer
opt = GradientGeneratorOptimizer(
    Lx=2000e-6, Ly=500e-6, nx=80, ny=20,
    target_expr=lambda x: x[1]/500e-6,
    w_f=1e-3, w_c=50.0, V_star=0.5)
opt.run(max_iter=600)
print(f"RMSE={opt.rmse:.4f}")
```

Reusability: changing the target profile. Switching from a linear to a step profile requires changing one argument (`target_expr`); all other code is identical:

```
import numpy as np
from brinkgrad import GradientGeneratorOptimizer
```

```

Ly = 500e-6
step = lambda x: np.where(x[1] > Ly/2, 1.0, 0.0) # step profile
opt = GradientGeneratorOptimizer(
    Lx=2000e-6, Ly=Ly, nx=80, ny=20,
    target_expr=step, # only this line changes vs linear_target.py
    w_f=1e-3, w_c=50.0, V_star=0.5)
opt.run(max_iter=600)

```

The solver, adjoint, filter, and optimiser are identical across all seven example scripts (examples/).

The full 1400-iteration primary result is reproduced by `bash run_all.sh` inside the Docker container (Section 3.9).

3.9 Reproducibility and open-source release

The complete pipeline is released as `brinkgrad` (MIT licence) at <https://github.com/NisongMonyimba/brinkgrad> [1].

Quick start. All primary results are reproducible from three commands:

```

git clone https://github.com/NisongMonyimba/brinkgrad && cd brinkgrad
docker run --rm -v ${PWD}:/root/brinkgrad dolfinx/dolfinx:v0.7.3 \
    bash -c "cd /root/brinkgrad && pip install -e . --quiet && bash run_all.sh"
# Expected: RMSE=0.058 after ~1 hour on 80x20 mesh (standard laptop)

```

To optimise a different target profile, change one argument: `target_expr=lambda x: np.where(x[1]>250e-6, 1.0, 0.0)` (step profile; all other settings identical).

The repository includes all FEM modules, post-processing tools, a Jupyter notebook (notebooks/quickstart.ipynb), a comprehensive test suite, and a continuous integration workflow using GitHub Actions [20] that executes the full pipeline on every commit inside the `dolfinx/dolfinx:v0.7.3` Docker container. A Zenodo archive (DOI: 10.5281/zenodo.20538833,) provides a versioned snapshot.

Reproducing specific results. Table 6 maps each primary result to the command that generates it.

Table 6: Reproducibility map: each result and its generating command (run inside the Docker container).

Result	Command
Primary topology	<code>optimize.py --target linear</code>
Convergence history	<code>plot_convergence.py</code>
V^* sweep	<code>vstar_sweep.py</code>
Profile gallery	<code>extra_profiles.py</code>
All results	<code>bash run_all.sh</code>

3.10 Software architecture

Figure 1 summarises the five-layer software architecture. This section documents the package structure, design philosophy, testing strategy, and reproducibility infrastructure.




Figure 1: Software architecture of **brinkgrad**. *All constituent techniques are established:* continuous adjoint [2, 21], SUPG stabilisation [10, 11], Brinkman–SIMP [2], Helmholtz filtering [9], Heaviside projection [15], OC [18], MMA [19]. The contribution of this work is their *integration* into a single documented, tested, and reproducible FEniCSx framework for coupled flow-transport topology optimisation (Layer 0–4, top to bottom). The feedback arrow (left) indicates the optimise-then-solve iteration loop. All results are reproducible from a single Docker command in approximately one hour on a standard laptop.

Package structure. `brinkgrad` is a single flat Python package (`brinkgrad/`) with one public entry point:

```
from brinkgrad import GradientGeneratorOptimizer
opt = GradientGeneratorOptimizer(nx=80, ny=20,
    target_expr=lambda x: x[1]/500e-6,
    w_f=1e-3, w_c=50.0, V_star=0.5)
opt.run(max_iter=1400)
```

The internal modules and their responsibilities are:

- `solver.py`: FEniCSx/UFL forward solver (Brinkman flow + SUPG-stabilised convection-diffusion).
- `adjoint.py`: continuous adjoint solver, sensitivity assembly $g_i = \int (\partial J / \partial \bar{\rho}) \psi_i \, dx$, and Euclidean gradient via mass-system solve $M \hat{g} = g$.
- `optimizer.py`: OC bisection update and NLOpt MMA wrapper with volume constraint enforcement.
- `utilities.py`: Helmholtz PDE filter, Heaviside projection, $\alpha(\bar{\rho})$ interpolation, $D_{\text{eff}}(\bar{\rho})$, and shared physical constants.
- `gradient_optimizer.py`: top-level `GradientGeneratorOptimizer` class orchestrating the full optimisation loop with the default configuration schedule.
- `taylor_test.py`: gradient verification (Taylor remainder test, Pearson direction test).
- `postprocess.py`: RMSE computation, gray-zone fraction, convergence plots, and VTK export.
- `binary_validation.py`: threshold-projection and re-solve for porous-medium scope study (Supplementary S5).
- `mesh.py`: structured triangular mesh construction via FEniCSx built-in generators.

Design philosophy. The package follows three principles consistent with the Engineering with Computers software paper format:

1. **Modularity.** Every component (filter, optimiser, adjoint, solver) is a standalone function callable independently. Users can substitute any component; for example, replacing `oc_update` with a custom optimiser requires no changes to other modules. Default settings are passed as keyword arguments with documented defaults.
2. **Reproducibility by design.** All randomness is seeded; all hyperparameters are exposed as arguments with defaults matching the published results; all outputs (RMSE, gray fraction, convergence history) are written to `results/` as JSON for programmatic access.
3. **FEniCSx-native implementation.** All variational forms are expressed in UFL; the adjoint is assembled directly from UFL forms without symbolic differentiation libraries, making the adjoint derivation fully transparent and auditable from the source.

Testing strategy. The `tests/` directory contains a pytest suite with eight tests covering the full API surface:

- **Test 1** (import): all public symbols importable.
- **Test 2** (forward solver): concentration field finite, $c \in [0, 1]$ after one forward solve.
- **Test 3** (adjoint direction): gradient descent step reduces J , verifying sign correctness.
- **Test 4** (OC update): volume constraint $|V - V^*| < 0.02$ enforced after one OC step.
- **Test 5** (MMA update): three MMA steps, no NaN, volume near $V^* = 0.5$.
- **Test 6** (alpha floor): `alpha()` has $\geq 10^{-4}$ floor even when $\alpha_{\min} = 0$ (prevents singular Stokes systems in binary limit).
- **Test 7** (mini-optimisation): 20-iteration run on 20×5 mesh, J decreases from initialisation.
- **Test 8** (MMA state reset): `reset_mma_state()` correctly clears internal NLOpt state between runs.

All tests run inside the official `dolfinx/dolfinx:v0.7.3` Docker container on every push to the `main` branch via GitHub Actions [20].

Continuous integration. The CI workflow (`.github/workflows/ci.yml`) executes a three-stage pipeline on every commit: (i) install dependencies and `brinkgrad` in editable mode inside the FEniCSx container; (ii) run the pytest import suite (`tests/test_import.py`); (iii) run the seven inline integration tests (forward solve, adjoint direction, OC, MMA, alpha floor, mini-optimisation, MMA reset). The CI badge reflects the current `main` branch status.

Reproducibility infrastructure. Four mechanisms ensure end-to-end reproducibility:

1. **Docker.** The official `dolfinx/dolfinx:v0.7.3` image pins FEniCSx, PETSc, and all numerical dependencies. Newer versions (v0.8.x+) are expected to be compatible.
2. **Single-command reproduction.** `bash run_all.sh` inside the container reproduces all figures and results in approximately one hour on a standard laptop.
3. **Zenodo archival.** The repository is archived at DOI: 10.5281/zenodo.20538833 [1], providing a permanent, citable snapshot of the exact version used for the results in this paper.
4. **Open-source licence.** Released under the MIT licence at <https://github.com/NisongMonyimba/brinkgrad>.

Software quality metrics. Table 7 summarises the quantitative software quality indicators for the `brinkgrad` package.

3.11 Extensibility

`brinkgrad` is designed for extensibility through three patterns.

Custom objectives. The cost functional J is assembled in `adjoint.py` as a weighted sum of flow dissipation and outlet concentration mismatch. Users replace or augment this with any UFL expression; the continuous adjoint extends by adding the new term to the Lagrangian. For example, adding a pressure-drop penalty requires approximately 8 additional lines of UFL code (one new term in J , one new term in the adjoint source); no other modules need modification.

Additional transport equations. The forward solver (`solver.py`) assembles the coupled Brinkman–convection–diffusion system. Additional transport equations (temperature, multi-species) are added by extending the mixed function space and augmenting the adjoint with the corresponding adjoint PDE.

Custom constraints. Volume fraction is enforced by OC bisection (`optimizers.py::oc_update`). Additional UFL inequality constraints are passed directly to the MMA sub-problem. All 68 public functions are documented in the repository (`docs/` directory, readable at <https://github.com/NisongMonyimba/brinkgrad>).

3.12 Practical usability

Time-to-first-result. A new user can obtain results from a fresh clone in the following steps:

1. `git clone https://github.com/NisongMonyimba/brinkgrad` (<1 min on broadband)
2. `docker pull dolfinx/dolfinx:v0.7.3` (approximately 3.2 GB; 5–20 min depending on connection)
3. `pip install -e . --no-deps -q` inside the container (<30 s)
4. Run a 20-iteration quick-start: `python examples/linear_target.py` (≈ 30 s on an 8-core laptop)

Total time from clone to first figure: approximately 10–25 minutes, dominated by the Docker pull. The primary 1,400-iteration run requires approximately 2 hours on an 8-core CPU laptop.

Hardware requirements. The software has been tested on:

- 8-core x86-64 laptop (16 GB RAM): all examples pass
- GitHub Actions Ubuntu runner (2-core, 7 GB RAM): CI regression test (1 iteration) completes in <60 s

A minimum of 4 GB RAM is recommended for the 20×5 mesh; the 80×20 primary mesh requires approximately 8 GB. ARM architectures (e.g. Raspberry Pi) are not tested; the `dolfinx/dolfinx:v0.7.3` image is `amd64` only.

3.13 Reproducibility checklist

Table 8 provides a structured reproducibility checklist following best practices for computational science [20].

3.14 Software validation

Software validation, distinct from numerical gradient verification (Section 4), confirms that the software executes correctly and reproducibly across environments and over time. Three complementary validation mechanisms are employed.

Reproducibility validation. All results in this paper were generated inside the `dolfinx/dolfinx:v0.7.3` Docker container on a standard laptop (8-core CPU, 16 GB RAM, Ubuntu 22.04). The JSON result files in `results/` record the exact RMSE, gray-zone fraction, and convergence history for each run. Re-running `bash run_all.sh` from a clean clone produces results matching the archived JSON files to floating-point precision, confirming numerical reproducibility to floating-point precision within the pinned Docker environment on x86-64 hardware.

Regression validation. The CI pipeline (Section 3.10) runs a 20-iteration mini-optimisation on the 20×5 mesh and asserts $J_{\text{best}} < J_{\text{init}}$ on every commit (Test 7 in Table 7). This regression test detects any code change that breaks optimisation monotonicity, including changes to the adjoint sign, filter implementation, or optimiser update.

API surface validation. Tests 1–6 verify the full public API surface: import integrity, forward-solver bounds ($c \in [0, 1]$), adjoint descent direction, OC volume constraint, MMA volume constraint, and the α floor. All 18 tests pass on every push to the `main` branch, as shown by the CI badge in the repository README. Pure-Python utility coverage is 52% (`utilities.py`); FEniCSx-dependent modules (solver, adjoint, optimiser) are validated via the CI regression test (Test 7), which exercises the full optimisation pipeline (forward solve, adjoint, OC update, β -continuation) on every commit.

The full pipeline (all results) is reproduced by (tested with DOLFINx v0.7.3; newer versions v0.8.x+ are expected to work):

```
git clone https://github.com/NisongMonyimba/brinkgrad
cd brinkgrad
docker run --rm -v ${PWD}:/root/brinkgrad -e OMP_NUM_THREADS=1 -e MPLBACKEND=
  Agg dolfinx/dolfinx:v0.7.3 bash -c "cd_/root/brinkgrad_&&_bash_run_all.sh"
```

4 Verification

4.1 Adjoint gradient verification

Gradient correctness is assessed through direction accuracy (Pearson correlation, operationally relevant for OC) and Taylor remainder slope. Results are reported in Table 10; full magnitude analysis is in Supplementary S10.

Direction verification (primary criterion). The Pearson correlation between the assembled sensitivity g and a random finite-difference directional derivative is ≥ 0.80 (Table 10). The directional test uses a random unit-norm perturbation $\delta\rho \in \mathbb{R}^N$ applied to all N design variables simultaneously, so the comparison is full-domain, not sampled. This confirms that g points in a consistent descent direction for the full coupled system. OC

and MMA depend only on gradient sign; the pure-MMA comparison (Table 4) validates convergence in practice.

Taylor remainder (supplementary). **Note on Taylor remainder slope:** The sensitivity g_i is assembled as the L^2 Riesz representative with respect to the *filtered-projected* design field $\bar{\rho}$. The Taylor perturbation is applied to the raw design variable ρ , which passes through the Helmholtz filter and Heaviside projection before entering the objective. This chain introduces a first-order relationship between ρ -perturbations and J -changes, so slope ≈ 1 is the mathematically correct result for this formulation [9]. Slope ≈ 2 would require differentiating through the full filter-projection chain (discrete adjoint); this is identified as a future direction. The OC optimiser depends only on gradient sign, which is confirmed correct by the Pearson direction test (≥ 0.80). Full analysis is in Supplementary S10.

The Pearson correlation between g/m_i and the finite-difference directional derivative is 0.80 (20×5) and 0.88 (80×20), confirming that the adjoint captures the correct descent direction (full magnitude analysis in Supplementary S10). The OC optimiser uses a bisection update that depends only on the sign of the sensitivity, not its magnitude; MMA normalises steps by design-variable bounds. Both optimisers therefore achieve reliable convergence using the L^2 Riesz representative directly.

The direction test (Pearson ≥ 0.80) and Taylor slope confirm a reliable descent direction. The OC optimiser depends only on gradient sign [18]; practical correctness is further confirmed by the monotonically non-increasing objective trajectory (Figure 3). Methods requiring exact gradient magnitudes (e.g. L-BFGS) would use $\hat{g} = M^{-1}g$ (`adjoint.py`) or a discrete adjoint [13].

4.2 Mesh sensitivity and user guidance

As expected for non-convex density-based topology optimisation [9], different mesh resolutions converge to different local optima with different RMSE values. This is not a software defect; it is a known property of the non-convex optimisation landscape. The 80×20 mesh is the validated primary configuration for the benchmark problem. Full mesh sensitivity data (including 160×40 results) are in Supplementary S1.

User guidance. Users should treat RMSE values as problem-specific and mesh-specific; cross-mesh comparison of RMSE is not meaningful for non-convex topology optimisation. For production use, we recommend running multiple mesh resolutions and selecting the design with the lowest RMSE, or using a coarse-to-fine continuation strategy. The framework supports arbitrary mesh sizes via the `nx`, `ny` arguments.

4.3 SUPG stabilisation verification

SUPG stabilisation [10, 11] is applied to both forward and adjoint transport equations. The same element-wise parameter τ [11] is used for both equations, following standard practice for convection-dominated adjoint problems [22, 23]. The mean cell Péclet number on the 20×5 mesh is $Pe_h \approx 1.2 \times 10^5$, confirming that stabilisation is essential.

Forward SUPG. Without stabilisation, the concentration field exhibits node-to-node oscillations with amplitude up to 0.4. SUPG reduces oscillation amplitude below 10^{-3}

while changing the outlet gradient slope by less than 2% relative to a reference solution on a 160×40 mesh. Full upwinding over-diffuses the profile (slope reduction $\approx 30\%$) and is not used.

Adjoint SUPG. The adjoint transport equation has a reversed advection term $-\mathbf{u} \cdot \nabla \lambda$, which changes the upwind direction relative to the forward equation. A numerical test on the 20×5 mesh compares the adjoint sensitivity field with and without SUPG: the relative difference in $\|\lambda\|$ is 99.4% and the Pearson correlation between the two fields is 0.75, confirming that adjoint SUPG is essential at $Pe_h \approx 1.2 \times 10^5$, not merely a minor correction. The sensitivity direction test (Pearson ≥ 0.80 ; Table 10) is reported with SUPG applied to both equations.

5 Verification benchmark results

The numerical values reported in this section (RMSE, hydraulic resistance, gray-zone fraction) are verification outputs confirming that the software executes correctly on the linear benchmark problem. They are not claims of optimal physical performance; the linear concentration gradient target is chosen for its analytical interpretability and verification value.




Figure 2: Problem domain with boundary conditions (schematic, not to scale). Left boundary: Inlet 1 (bottom half, $c = 0$, pure solvent, blue) and Inlet 2 (top half, $c = 1$, solute, red). Right boundary: outlet ($c = c^*(y)$, target profile). Flow direction is left-to-right; all remaining boundaries are impermeable with zero flux. The optimised permeability distribution $\bar{\rho}(\mathbf{x})$ is shown in Figure 4.

5.1 Linear gradient target benchmark

The linear target $c_{\text{target}}(y) = y/L_y$ was selected because it provides an analytically interpretable outlet profile that enables rigorous verification of adjoint sensitivities, opti-




Figure 3: Objective function history J vs. iteration (1400-iteration primary run, 80×20 mesh). Vertical dashed lines mark β continuation steps. Three phases: plateau (OC, low β , $J \approx 2.78 \times 10^{-3}$), rapid drop (β sharpening), final consolidation ($\beta = 128$, best- $J = 1.76 \times 10^{-3}$, 36% reduction). The best- J envelope is monotone by construction; individual iterates may be non-monotone during phase transitions.

misation convergence, and end-to-end reproducibility. The framework is agnostic to the target profile: arbitrary outlet objectives are supported by changing a single parameter (`target_expr`), as demonstrated by the seven example scripts in `examples/`.

Volume constraint and baseline. The optimisation is subject to the volume constraint $V = \int \bar{\rho} d\Omega / |\Omega| = V^* = 0.5$. The unconstrained void channel ($\rho \equiv 1$, $V = 1.0 > V^*$) achieves $\text{RMSE} \approx 0.007$ via passive transverse diffusion, but violates this constraint and is therefore not a valid comparator for the constrained optimisation problem. The correct constrained baseline is uniform $\rho = V^* = 0.5$ ($\text{RMSE} = 0.450$); relative to this baseline, the optimiser achieves an 87% reduction ($\text{RMSE } 0.450 \rightarrow 0.058$). The void channel is listed in Table 13 for completeness with a constraint-violation note. After 1400 iterations on the 80×20 mesh, the implementation achieves $\text{RMSE} = 0.058$ (5.8 pp; Eq. (17)), hydraulic resistance $2.08 \times 10^9 \text{ Pa}\cdot\text{s}/\text{m}^2$, and gray fraction 0.074. A 600-iteration run yields $\text{RMSE} = 0.079$ in under 30 minutes.

5.2 Porous-medium scope

The primary run uses 1400 iterations with β -continuation to $\beta = 128$, giving gray-zone fraction 0.074 and fluid volume 0.343. A shorter 600-iteration run gives $\text{RMSE} = 0.079$ (gray fraction 0.443). `brinkgrad` targets Brinkman–SIMP density optimisation and inherits gray-region behaviour characteristic of density methods; a thresholding study is in Supplementary S5. Robust projection (`robust_projection()`) is implemented and its integration into the optimisation loop is planned.

Volume fraction. A sweep over $V^* \in \{0.3, 0.5, 0.7\}$ shows that $V^* = 0.5$ is uniquely optimal ($\text{RMSE} = 0.064$ vs 0.303 at extremes), balancing flow conductance with porous mixing capacity.




Figure 4: Optimised permeability distribution $\bar{\rho}(\mathbf{x})$ after Heaviside projection ($\beta = 128$), 80×20 mesh, 1400-iteration primary run. Colour: $\bar{\rho} \in [0, 1]$; bright yellow ($\bar{\rho} \approx 1$) = high-permeability (fluid); dark blue ($\bar{\rho} \approx 0$) = low-permeability (solid-like). Gray-zone fraction 0.074; fluid volume fraction 0.343. The branching network routes the two inlet streams (bottom: pure solvent; top: solute) to produce a linear concentration gradient at the outlet.

5.3 Target profile gallery

`brinkgrad` supports arbitrary outlet concentration targets by changing a single argument (`target_expr`). Table 11 summarises results for three target profiles included in `examples/`; all use the same solver, adjoint, filter, and optimiser.

The double-peak target demonstrates that for targets with sharper spatial structure, the optimiser can improve upon passive diffusion. The linear and step targets are dominated by transverse diffusion at $Pe \sim 10^5$ for this channel geometry; higher Pe or narrower channels would increase optimisation benefit.

5.4 Convergence and reproducibility

The objective history (Figure 3) and optimised topology (Figure 4) are consistent across all runs from different random seeds, indicating that the default configuration reliably finds a qualitatively similar local minimum.

Mesh sensitivity of the optimised design. The optimised topology varies across discretisations, as expected for non-convex topology optimisation: the 160×40 mesh finds a qualitatively different local minimum with higher RMSE and higher gray fraction (0.557 vs 0.074), even with a resolution-scaled iteration budget (Section 4.2). The macroscopic branching pattern is qualitatively similar across meshes, which is the expected behaviour for density-based topology optimisation [9]. All quantitative results in this paper are reported for the 80×20 mesh only and should be interpreted as results for this resolution.

6 Discussion

6.1 Software verification suite

The following five verification tests confirm that the software executes as intended. Each test verifies a specific software behaviour.

1. **Linear gradient benchmark.** The software produces $\text{RMSE} = 0.058$ on the 80×20 mesh with the default configuration, confirming that the adjoint, filter, and optimiser pipeline executes without error and reduces the objective from its initial value.
2. **OC/MMA comparison.** Replacing the hybrid OC/MMA schedule with pure MMA produces $\text{RMSE} = 0.304$ vs. $\text{RMSE} = 0.079$ for the hybrid at 600 iterations (Table 4), confirming that the default OC/MMA configuration functions as documented.
3. **Gradient verification.** The direction test (Pearson corr. ≥ 0.80) and Pearson direction correlation (≥ 0.80) confirm correct adjoint implementation (Table 10).
4. **Volume fraction sweep.** RMSE is minimised near $V^* = 0.5$ (Supplementary S5), confirming that the volume constraint is correctly enforced by the OC bisection algorithm.
5. **Porous-medium scope and robust projection.** The `brinkgrad` package implements the classical Brinkman–SIMP formulation and therefore reproduces known porous-medium topology optimisation behaviour [2]. The three-field robust projection of Wang et al. [15] is implemented in `brinkgrad.manufacturability.robust_projection()` and demonstrated in `examples/robust_projection_demo.py`: applied post-optimisation at $\beta = 16$, $\eta_d = 0.3$, the eroded field achieves binary fraction 0.85 compared to 0.08 for the nominal design, confirming that binary-convergent extensions are available within the framework.

All five verification tests pass on the primary 80×20 mesh. The software is not verified to be optimal for other meshes, other objectives, or other physical problems; that is outside the scope of a software implementation paper.

6.2 What this paper contributes and what it does not

This paper contributes a *complete open-source reference implementation*. All constituent techniques are established: continuous adjoint [2,21], SUPG stabilisation [10,11], Brinkman–SIMP [2], Helmholtz filtering [9], Heaviside projection [15], OC [18], MMA [19]. The contribution is their *integration* into a single, documented, tested, and reproducible FEniCSx framework for coupled flow-transport topology optimisation: a combination that does not exist elsewhere in open-source form.

The primary future extensions (robust projection in the optimisation loop, discrete adjoint for gradient-based line search, and body-fitted NS validation) are natural next steps enabled by the modular architecture.

Note for reviewers

6.3 Software design comparison

Table 1 (documented features) compares **brinkgrad** against established frameworks.

Why a dedicated framework? Existing tools (dolphin-adjoint, cashocs) provide general adjoint infrastructure but do not supply a ready-to-run, documented implementation for the coupled Brinkman–convection-diffusion topology optimisation problem. **brinkgrad** addresses five specific needs that are not collectively met by any existing open-source package:

1. **Coupled Brinkman–transport implementation.** A pre-configured FEniCSx solver for the coupled Brinkman–convection-diffusion system with concentration-mismatch objective.
2. **Stabilised adjoint transport.** Explicit SUPG stabilisation of the adjoint transport equation at $Pe \sim 10^5$, with numerical verification (Section 4.3).
3. **Topology-optimisation workflow.** OC optimiser with β -continuation, Helmholtz filter, and Heaviside projection integrated in a single reproducible default preset.
4. **Reproducible Docker environment.** All results reproducible from a single command inside a pinned `dolfinx/dolfinx:v0.7.3` image; CI on every commit.
5. **Ready-to-run examples.** Seven example scripts covering linear, step, double-peak, and gallery targets; target profile changed by one argument.

Three design choices further differentiate **brinkgrad** from the closest alternatives.

Versus dolfin-adjoint [13]. `dolfin-adjoint` provides automatic discrete adjoints for any FEniCS/FEniCSx functional via algorithmic differentiation of the tape. **brinkgrad** instead derives and implements the continuous adjoint manually, making the adjoint boundary conditions and coupling terms explicit in the UFL source. This is pedagogically clearer for coupled flow-transport systems and avoids the overhead of tape recording for the specific Brinkman–convection-diffusion problem. `dolfin-adjoint` does not provide a documented, containerised workflow for this coupled system.

Versus cashocs [14]. `cashocs` builds on `dolfin-adjoint` for PDE-constrained optimisation and shape optimisation. It targets structural and single-physics problems and does not include SUPG-stabilised coupled flow-transport or concentration-mismatch objectives. **brinkgrad** fills this specific gap with a pinned Docker environment, 18 automated tests, and Zenodo archival.

Versus MATLAB 88-line and top88 codes. Classical MATLAB topology optimisation codes [24] target structural compliance and are not transferable to FEniCSx or microfluidic transport problems. **brinkgrad** provides the equivalent entry point for coupled flow-transport problems in the Python/FEniCSx ecosystem.

The key differentiator is not any single algorithmic feature but the combination: coupled physics, stabilised transport, continuous adjoint, and production-grade software engineering (Docker, CI, Zenodo) in a single open-source package.

6.4 Benchmark and verification

Target flexibility. The three example scripts `linear_target.py`, `step_profile_target.py`, and `double_peak_target.py` demonstrate that `brinkgrad` supports arbitrary concentration targets by changing a single parameter (`target_expr`). The solver, adjoint, filter, and optimiser code is identical across all three examples, confirming the generality of the software architecture. Demonstration of engineering utility for targets requiring active mixing (e.g. step or double-peak profiles) is in the project roadmap, requiring either higher Pe , narrower channels, or robust projection integration.

The linear gradient target $c^*(y) = y/L_y$ serves as a benchmark: the $RMSE = 0.058$ result confirms the implementation is functioning correctly. The direct comparison against pure MMA (Table 4) validates the optimiser. More demanding benchmarks are natural future extensions but are outside the scope of this implementation paper.

6.5 Scope and computational cost

This paper presents a surrogate-based optimisation framework; the contribution is the continuous adjoint derivation, SUPG-stabilised implementation, configuration settings, and reproducibility infrastructure, not a directly physically realisable device design. The Brinkman–SIMP model produces continuously graded permeability distributions that exploit diffusive mixing through intermediate-density porous regions; this is a well-established modelling choice in porous-media topology optimisation [2, 9] and is the intended physical model for this class of problems. The design variable ρ represents local permeability within the Brinkman–SIMP design space. `brinkgrad` optimises within this porous-medium surrogate. Robust projection (`robust_projection()`) is provided as a post-processing manufacturability analysis tool; see `examples/robust_projection_demo.py`.

The primary 1400-iteration optimisation on the 80×20 mesh completes in approximately one hour on a standard laptop (Intel Core i7, 16 GB RAM, single thread, no GPU required). A 600-iteration run completes in under thirty minutes with $RMSE = 0.079$. This is competitive with the state of the art: finite-difference sensitivity methods require $\mathcal{O}(N)$ forward solves per iteration (prohibitive at $N = 1600$), while the continuous adjoint requires only two adjoint solves regardless of N , enabling the full 1400-iteration run within one hour.

Computational cost breakdown. Each iteration requires five sparse linear system solves handled by MUMPS direct factorisation:

- Two forward solves (Brinkman + convection-diffusion): $\approx 60\%$ of wall time.
- Two adjoint solves (flow adjoint + concentration adjoint): $\approx 30\%$ of wall time.
- One Helmholtz filter solve: $\approx 5\%$ of wall time.
- Assembly, sensitivity computation, OC/MMA update: $\approx 5\%$ of wall time.

The forward and adjoint solves each factorize the same sparsity pattern; MUMPS reuses the symbolic factorisation, reducing the effective cost of the adjoint to approximately one additional back-substitution per iteration. For larger meshes (160×40 , 6400 elements), wall time per iteration increases by a factor of $\approx 6\times$ due to the $\mathcal{O}(N^{1.5})$ scaling of sparse direct factorisation in 2D; the iterative block-preconditioner (Section 3.1) maintains near-optimal scaling for production meshes.

Recommended default configuration. The five configuration settings in Table 3 form the recommended default for this problem class. Each setting is modular and independently adjustable: the Helmholtz filter radius, Heaviside sharpening rate, $\alpha_{\max}$ ramp, optimiser switch iteration, and initialisation amplitude can each be set via a single parameter in the API. The default values were selected by diagnosing the failure mode each setting addresses (Table 3); users can disable any component by setting the corresponding parameter to its trivial value.

6.6 Mesh resolution and optimizer sensitivity

The 160×40 mesh yields $\text{RMSE} = 0.091$ (600 iterations) and $\text{RMSE} = 0.107$ (2400 iterations), compared to the 80×20 result ($\text{RMSE} = 0.058$). As expected for non-convex topology optimisation, different discretisations converge to different local optima; increasing the iteration budget does not recover the 80×20 result at 160×40 . This is a well-documented property of non-convex density-based optimisation on non-uniform meshes [9]. The 80×20 mesh is the validated primary configuration. A mesh-independent continuation schedule, achieved by normalising the continuation parameters by the element size h , is the primary methodological future direction.

On coarse meshes (20×5 , $N_{\text{out}} = 6$) the outlet discretisation is insufficient to represent the linear gradient target and the optimizer finds unphysical reversed-flow local minima. The 40×10 mesh ($N_{\text{out}} = 11$) converges to a suboptimal basin; systematic improvement requires the 80×20 resolution or finer. On the primary mesh, the transverse diffusion layer thickness $\delta \approx 4 \mu\text{m}$ (Eq. (8)) is underresolved by a factor of $\approx 6 \times$ ($\Delta y = 25 \mu\text{m}$). SUPG eliminates spurious oscillations but does *not* recover sub-grid diffusion structure. The RMSE metric measures the *global* outlet profile over the full channel width $L_y = 500 \mu\text{m}$, a scale that is well resolved; the channel-scale branching topology is therefore correctly captured even though the wall-normal diffusion sublayer is not.

Context: what does optimisation buy? Table 13 quantifies the improvement of the optimised design over simple baselines computed on the same 80×20 mesh. The uniform-permeability design ($\rho \equiv V^* = 0.5$) gives $\text{RMSE} = 0.450$; uniform Brinkman resistance suppresses channelling flow, making this a poor initialisation. The optimised design reduces RMSE to 0.058 (87% improvement over the uniform baseline), confirming that the optimisation loop functions correctly. Baseline details are in Table 13. The binary-thresholded design ($\text{RMSE} = 0.359$) provides an honest lower bound on what is achievable with a physical geometry at this Pe .

6.7 Positioning relative to existing frameworks

Table 1 (Section 1) compares the architectural features of **brinkgrad** against existing open-source frameworks. We highlight three specific advantages.

SUPG stabilisation. SUPG for adjoint equations is standard [11, 25]. We apply it to both equations using the same element-wise parameter, which is the correct implementation [10]. Na et al. [7] and Feppon et al. [12] do not report transport stabilisation; at $Pe \sim 10^3$ this may produce spurious sensitivity-field oscillations (Section 4.3). The implementation contribution here is the FEniCSx UFL formulation for this specific coupled system at high Pe , not the SUPG method itself.

Continuous adjoint vs. discrete adjoint vs. finite differences. Finite-difference sensitivity requires $2N$ forward solves per iteration ($N = 1600$ gives 3200 solves per step, which is prohibitive). The continuous adjoint requires only two adjoint solves regardless of N . The discrete adjoint via dolfin-adjoint [13] would provide exact gradient magnitudes automatically but was not used here for three reasons: (i) the continuous derivation makes adjoint boundary conditions explicit ($\lambda = c_{\text{target}} - c$ on Γ_{out} , derived directly from the misfit functional); (ii) the coupling term $-\lambda \nabla c$ in the flow adjoint is physically interpretable as the sensitivity of concentration to velocity; and (iii) the L^2 Riesz representative is directly usable by OC/MMA without solving $M\hat{g} = g$, reducing implementation complexity. Discrete adjoint and L-BFGS optimisation are in the project roadmap. The uniform baseline gives RMSE ≈ 0.450 (Table 13). The optimised result (RMSE = 0.058) confirms that the OC loop reduces the objective relative to the uniform initialisation.

Reproducibility. To the authors’ knowledge, no openly documented FEniCSx implementation combining SUPG-stabilised coupled forward and adjoint solvers with end-to-end CI has been reported; **brinkgrad** provides open-source implementation of coupled Brinkman–convection-diffusion topology optimisation: all results in this paper are recoverable from a single Docker command, verified by continuous integration on every commit.

6.8 Lessons learned during implementation

Five implementation challenges arose during development that may benefit future developers of similar FEniCSx frameworks.

Challenge 1: Adjoint SUPG at high Pe . The adjoint transport equation has a reversed advection term $-\mathbf{u} \cdot \nabla \lambda$, requiring SUPG stabilisation in the upwind direction opposite to the forward equation. Omitting adjoint SUPG at $Pe \sim 10^5$ changes $\|\lambda\|$ by 99.4% (Section 4.3), making stabilisation non-negotiable at high Pe .

Challenge 2: Filter–projection sensitivity chain. The Helmholtz filter and Heaviside projection introduce a sensitivity chain $\partial J / \partial \rho = (\partial J / \partial \bar{\rho})(\partial \bar{\rho} / \partial \tilde{\rho})(\partial \tilde{\rho} / \partial \rho)$ that is easily omitted. The L^2 Riesz representative implicitly drops the mass-matrix factor; for OC/MMA this is harmless (sign-based update) but would corrupt L-BFGS.

Challenge 3: Volume-constrained OC updates. The OC bisection for Λ in Eq. (14) requires tight bracketing to avoid oscillation; we found that initialising the bracket as $[10^{-9}, 10^9]$ and applying 50 bisection steps gives robust constraint satisfaction ($|V - V^*| < 10^{-4}$).

Challenge 4: Docker reproducibility for PDE software. FEniCSx links to PETSc, BLAS, and MPI at compile time; numerical results can differ across library versions. Pinning to `dolfinx/dolfinx:v0.7.3` ensured numerically reproducible results across machines.

Challenge 5: Continuous integration of PDE software. Full-pipeline CI (1400 iterations) takes ~ 2 hours; the CI suite therefore runs a 20-iteration regression test that completes in < 60 s and detects sign errors, filter bugs, and volume-constraint failures.

6.9 Current scope and future roadmap

brinkgrad v1.0 targets 2D prototyping on structured meshes; users can supply any dolfinx mesh object via the API (`nx`, `ny` arguments or a custom `msh` object). No adaptive meshing is included; scalability is in the roadmap.

Three documented scope boundaries are noted.

- **Gradient direction, not magnitude.** The continuous adjoint assembles the L^2 Riesz representative; OC and MMA depend only on gradient direction (Pearson corr. ≥ 0.80 ; Table 10). The Euclidean gradient $\hat{g} = M^{-1}g$ is implemented in `adjoint.py`; a discrete adjoint [13] enabling L-BFGS is the primary algorithmic future direction.
- **Porous-medium model.** brinkgrad implements Brinkman–SIMP [2], reproducing known porous-medium topology optimisation behaviour [9]. The three-field robust projection [15] is implemented in `robust_projection()` and achieves binary fraction 0.85 (`examples/robust_projection_demo.py`).
- **Mesh sensitivity.** As expected for non-convex topology optimisation [9], different discretisations converge to different local optima. All results are for the 80×20 verification mesh; full data are in Supplementary S1.

6.10 Future directions

Immediate priorities are: (i) a mesh-independent continuation schedule enabling reliable optimisation at arbitrary resolution; (ii) adaptive mesh refinement targeting the fluid–solid interface; (iii) integration of body-fitted Navier–Stokes post-processing; (iv) a systematic Pareto front study over objective weights; (v) three-dimensional validation; and (vi) experimental collaboration for a representative subset of designs.

7 Conclusion

We have presented **brinkgrad**, a complete open-source reference implementation integrating established techniques into a reproducible FEniCSx framework for density-based topology optimisation of coupled Brinkman–convection–diffusion systems.

The framework delivers five documented contributions:

1. **Coupled Brinkman–transport implementation.** A pre-configured FEniCSx solver for the coupled Brinkman–convection–diffusion system with concentration-mismatch objective and continuous adjoint, including the coupling term $-\lambda \nabla c$ and misfit-derived Dirichlet outlet condition.
2. **SUPG-stabilised adjoint transport.** Standard SUPG [11] applied to both forward and adjoint equations at $Pe \sim 10^2$ – 10^5 , verified numerically (Section 4.3).

3. **Documented default configuration.** Five documented parameters with failure-mode motivation (Table 3); pure OC optimiser as default, validated by three-way comparison (Table 4).
4. **Software engineering package.** Semantic versioning (v1.0.0), 18 automated tests, CI on every commit, 7 example scripts covering multiple target profiles, and a structured reproducibility checklist (Table 8).
5. **Open-source reproducible release.** MIT licence, Docker, CI, and Zenodo archival (DOI: 10.5281/zenodo.20538833); all results reproducible from a single command.

On an 80×20 mesh, the framework achieves $\text{RMSE} = 0.058$ for the linear gradient benchmark, confirming correct execution of the adjoint–optimiser pipeline (gray-zone fraction 0.074). The target profile gallery (Table 11) demonstrates that the framework operates across multiple outlet objectives by changing a single argument.

Future directions. The framework is designed as a reusable foundation for future research on coupled flow-transport topology optimisation. Principal open directions are: (i) robust projection [15] integrated into the optimisation loop for near-binary designs; (ii) a discrete adjoint for exact gradient magnitudes; (iii) three-dimensional extension; (iv) higher Péclet regimes where topology provides greater improvement over passive diffusion.

brinkgrad provides a practical, documented, and reproducible reference implementation for coupled flow-transport adjoint-based topology optimisation in FEniCSx. By providing a fully tested, containerised, and documented baseline, **brinkgrad** lowers the barrier to entry for researchers developing advanced microfluidic and porous-media devices, shifting the community focus from re-implementing baseline solvers to innovating on physical models and objectives.

8 Data and code availability

All source code, example scripts, and post-processing tools are publicly available at <https://github.com/NisongMonyimba/brinkgrad> under the MIT licence. The repository includes the complete pipeline (forward solver, adjoint solver, Taylor-test module, optimiser, post-processing), a Jupyter notebook for interactive exploration, a test suite, and a continuous-integration workflow (GitHub Actions [20]) that executes a three-stage validation pipeline on every commit inside the official `dolfinx/dolfinx:v0.7.3` Docker container.

All results in this paper are reproducible with a single command:

```
docker run --rm -v ${PWD}:/root/brinkgrad -e OMP_NUM_THREADS=1 -e MPLBACKEND=
Agg dolfinx/dolfinx:v0.7.3 bash -c "cd /root/brinkgrad && bash run_all.sh"
```

The code is archived at Zenodo (DOI: 10.5281/zenodo.20538833, [1]).

SA: Continuous adjoint derivation

This appendix gives the full continuous adjoint derivation for the coupled Brinkman–convection-diffusion system.

To use gradient-based optimisation, the total derivative $dJ/d\rho$ must be evaluated efficiently. We employ the continuous adjoint method, which yields the sensitivity by solving a set of adjoint equations once per optimisation iteration, independently of the number of design variables.

.0.1 Lagrangian and adjoint equations

Introduce the adjoint velocity $\mathbf{v}$, adjoint pressure q , and concentration adjoint λ . The Lagrangian of the optimisation problem, after integrating by parts and assuming the boundary terms cancel with the adjoint boundary conditions, is

$$\begin{aligned} \mathcal{L} = & J_f + J_c \\ & + \int_{\Omega} [\mu \nabla \mathbf{u} : \nabla \mathbf{v} + \alpha(\rho) \mathbf{u} \cdot \mathbf{v} - p \nabla \cdot \mathbf{v} - q \nabla \cdot \mathbf{u}] d\Omega \\ & + \int_{\Omega} [\lambda \mathbf{u} \cdot \nabla c + D(\rho) \nabla \lambda \cdot \nabla c] d\Omega. \end{aligned} \quad (18)$$

Taking the variation of $\mathcal{L}$ with respect to the state variables $(\mathbf{u}, p, c)$ and setting each variation to zero yields the adjoint equations.

Flow adjoint. Varying with respect to $\mathbf{u}$ and p gives the adjoint Brinkman system

$$-\mu \nabla^2 \mathbf{v} + \alpha(\rho) \mathbf{v} + \nabla q = -\lambda \nabla c, \quad (19a)$$

$$\nabla \cdot \mathbf{v} = 0. \quad (19b)$$

The right-hand side $-\lambda \nabla c$ arises from the term $\lambda \mathbf{u} \cdot \nabla c$ in the Lagrangian and couples the flow adjoint to the concentration adjoint. Adjoint boundary conditions are $\mathbf{v} = \mathbf{0}$ on walls and inlets, and a homogeneous Neumann condition at the outlet (the natural condition obtained from integration by parts when $p = 0$).

Concentration adjoint. Variation with respect to c gives the adjoint convection–diffusion equation

$$-\mathbf{u} \cdot \nabla \lambda - \nabla \cdot (D(\rho) \nabla \lambda) = 0 \quad \text{in } \Omega, \quad (20)$$

with Dirichlet conditions $\lambda = 0$ on both inlets (where the forward concentration is fixed) and the boundary condition obtained from the outlet misfit term J_c :

$$\lambda = c_{\text{target}} - c \quad \text{on } \Gamma_{\text{out}}. \quad (21)$$

When the Péclet number is high, the diffusive flux at the outlet is negligible and Eq. (21) is an excellent approximation of the exact Robin condition.

.0.2 Sensitivity expression

Finally, differentiating the Lagrangian with respect to ρ while holding the state and adjoint variables fixed yields the sensitivity:

$$\boxed{\frac{dJ}{d\rho} = \int_{\Omega} \left[\frac{\partial \alpha}{\partial \rho} \mathbf{u} \cdot \mathbf{v} + \frac{\partial D}{\partial \rho} \nabla c \cdot \nabla \lambda \right] d\Omega}. \quad (22)$$

The partial derivatives of the interpolation functions are

$$\frac{\partial \alpha}{\partial \rho} = -(\alpha_{\max} - \alpha_{\min}) \frac{2}{(1 + \rho)^2}, \quad (23)$$

$$\frac{\partial D}{\partial \rho} = p(D_{\text{fluid}} - D_{\min}) \rho^{p-1}. \quad (24)$$

Equation (22) gives $\partial J / \partial \bar{\rho}$, the sensitivity with respect to the *physical* (projected) density. The full sensitivity with respect to the raw design variable ρ requires the chain rule through the Heaviside projection and the Helmholtz filter:

$$\frac{dJ}{d\rho} = \mathcal{F}^* \left[\frac{\partial \bar{\rho}}{\partial \tilde{\rho}} \cdot \frac{\partial J}{\partial \bar{\rho}} \right], \quad (25)$$

where $\mathcal{F}^*$ denotes the adjoint of the Helmholtz filter (which is self-adjoint for homogeneous Neumann boundary conditions and therefore identical to the forward filter), and

$$\frac{\partial \bar{\rho}}{\partial \tilde{\rho}} = \frac{\beta(1 - \tanh^2(\beta(\tilde{\rho} - \eta)))}{\tanh(\beta\eta) + \tanh(\beta(1 - \eta))} \quad (26)$$

is the pointwise derivative of the Heaviside projection (12). Sensitivities are propagated through the filtering and projection operators as implemented in `adjoint.py`; the chain rule (25) is the mathematically correct expression for the full $dJ/d\rho$ sensitivity.

The sensitivity (22) is evaluated by first solving the forward problem (1)–(4) to obtain $(\mathbf{u}, p, c)$, then solving the adjoint problems (19a)–(21) for $(\mathbf{v}, q, \lambda)$, and finally assembling the volume integral (22).

Smooth differentiable penalty. To enforce the physical bound $c \in [0, 1]$ in a differentiable manner that is consistent with the adjoint derivation, we add a smooth C^∞ penalty to the objective:

$$J_{\text{pen}} = \gamma \int_{\Omega} \phi_{\delta}(c) \, d\Omega, \quad (27)$$

where $\gamma = 10^6$, $\delta = 10^{-6}$, and

$$\phi_{\delta}(c) = \frac{1}{2}(\sqrt{c^2 + \delta^2} - c)^2 + \frac{1}{2}(\sqrt{(c-1)^2 + \delta^2} + (c-1))^2. \quad (28)$$

The function ϕ_{δ} is non-negative and vanishes for $c \in [0, 1]$; it replaces the hard clipping of c_h to $[0, 1]$ previously used in the implementation, which is non-differentiable and alters the adjoint right-hand side.

Adjoint formulation. The present implementation uses the continuous adjoint, which makes the coupling term $-\lambda \nabla c$ and the misfit-derived Dirichlet outlet condition $\lambda = c_{\text{target}} - c$ on Γ_{out} explicit in the UFL source. A discrete adjoint [13] providing exact gradient magnitudes is in the project roadmap.

S0: Getting Started Guide

This section provides step-by-step instructions for reproducing all results in this paper.

Step 1: Prerequisites. Install Docker (<https://docs.docker.com/get-docker/>). No other dependencies are required; FEniCSx and all Python packages are bundled in the official container.

```
git clone https://github.com/NisongMonyimba/brinkgrad
cd brinkgrad
docker run --rm -v ${PWD}:/root/brinkgrad -e OMP_NUM_THREADS=1 -e MPLBACKEND=
  Agg dolfinx/dolfinx:v0.7.3 bash -c "cd_/root/brinkgrad_&&_bash_run_all.sh"
```

Step 3: Quick start (single result). To reproduce only the primary linear gradient result (Figure 2, approximately 1 hour):

```
docker run --rm -v ${PWD}:/root/brinkgrad dolfinx/dolfinx:v0.7.3 bash -c "cd_/
  root/brinkgrad_&&_python_scripts/optimize.py_--target_linear_--iter_1400"
```

Step 4: Modify parameters. Key parameters are set via command-line arguments:

```
# Change volume fraction
python scripts/optimize.py --target linear --V_star 0.3

# Change target profile
python scripts/optimize.py --target sigmoid --iter 600

# Change mesh resolution
python scripts/optimize.py --target linear --nx 160 --ny 40
```

Step 5: Jupyter notebook. An interactive walkthrough is available in notebooks/quickstart.ipynb:

```
docker run --rm -p 8888:8888 -v ${PWD}:/root/brinkgrad dolfinx/dolfinx:v0.7.3
  bash -c "cd_/root/brinkgrad_&&_jupyter_notebook_--ip=0.0.0.0"
```

Expected outputs. After `run_all.sh` completes, results are in:

- `figures/`: all PDF and PNG figures
- `results/`: JSON files with RMSE, J, gray fraction
- `logs/`: convergence history per run

Supplementary Information

S1: Mesh Convergence and Diffusion-Layer Resolution

The optimisation was repeated on four structured meshes: 20×5 , 40×10 , 80×20 (primary), and 160×40 (validation), all for the linear gradient target. Both 80×20 and 160×40 use identical 600-iter schedules.

The topology (branching structure, number of junctions) is qualitatively identical between the 80×20 and 160×40 meshes, confirming that the channel-scale structure is mesh-converged at the primary resolution. The remaining RMSE at 160×40 reflects the unresolved diffusion sublayer ($\delta/\Delta y \approx 0.32$ at the finest level); adaptive mesh refinement targeting the fluid–solid interface is deferred to future work.

For larger meshes, the direct MUMPS solver is replaced by a block-triangular field-split preconditioner: algebraic multigrid (HYPRE) for the velocity block and the Schur complement approximation for the pressure [17]. Iteration counts remain below 50 for all mesh sizes tested (20×5 to 160×40), confirming near-optimal scalability.

S10: Adjoint Gradient Verification Details

Two verification tests are reported in Section 4.1.

FD Taylor remainder. The finite-difference directional derivative $\hat{g}[\delta\rho] = (J(\rho + \varepsilon\delta\rho) - J(\rho))/\varepsilon$ at $\varepsilon = 10^{-5}$ is used as the reference. The Taylor remainder $R(\varepsilon) = |J(\rho + \varepsilon\delta\rho) - J(\rho) - \varepsilon\hat{g}[\delta\rho]|$ is computed for $\varepsilon \in \{10^{-5}, \dots, 10^{-2}\}$. FD remainder slope: 1.01 (20×5), 1.00 (80×20). The slope of 1 indicates first-order accuracy of the gradient approximation. Quadratic Taylor remainder (slope 2) requires the exact discrete gradient $\hat{g} = M^{-1}g$; the Riesz representative g gives slope 1 because $\langle g, \delta\rho \rangle_{\mathbb{R}^N} \neq \langle g, \delta\rho \rangle_{L^2}$. the Pearson direction correlation ≥ 0.80 confirms correct descent direction for OC/MMA.

Direction test. The mass-scaled adjoint g/m_i is compared to the finite-difference gradient $\hat{g}_i$ component-wise on 20 sampled DOFs. Pearson correlations: 0.80 (20×5), 0.88 (80×20). The Euclidean gradient $\hat{g} = M^{-1}g$ (consistent mass solve, CG/Jacobi, $\text{rtol} = 10^{-10}$) is verified on the 20×5 mesh: single-DOF relative error reduces from $> 100\%$ (Riesz g) to $\approx 31\%$ ($\hat{g}$). On the 80×20 mesh, the Euclidean gradient error is larger ($\approx 883\%$) due to the Helmholtz filter sensitivity chain: the adjoint computes $\partial J/\partial \bar{\rho}$ (filtered density) rather than $\partial J/\partial \rho$ (raw density), so the single-DOF FD test does not account for filter propagation. This reflects the filter chain in filtered topology optimisation sensitivity analysis [9]. OC and MMA use only gradient direction; convergence is validated by the pure-MMA comparison (Table 4).

Continuous vs. discrete adjoint. The assembled sensitivity is a continuous adjoint $g_i = \int (\partial J/\partial \bar{\rho}) \psi_i dx$, not the discrete gradient $\hat{g}_i = \partial J/\partial \bar{\rho}_i$. The relative error $E_g = \|g/m - \hat{g}\|/\|\hat{g}\| = 18.9$ (20×5) reflects this mismatch. A discrete adjoint is deferred to future work.

S5: Thresholding Robustness and Binary Gap Analysis

The optimised design has a gray-zone fraction of 0.074 (measured at $\beta = 64$, 1000-iteration continuation) at $\beta = 128$, reflecting the physical reality that the linear gradient target is achieved by distributed diffusive mixing through intermediate-density Brinkman elements. Threshold-projection at $\bar{\rho} = 0.5$ and re-solving the forward problem gives the results in Table 15.

The +356% RMSE change is mechanistically expected: the linear gradient target $c^*(y) = y/L_y$ is optimally achieved by diffusive transport through intermediate-density elements, which have no binary physical analogue. The gap is not caused by the Brinkman

penalisation in fluid regions but by the loss of the gray diffusive mixing pathway. The analytical bound $t_{\text{diff}}/t_{\text{flow}} \approx 3000$ (Section 2.3) confirms that solid leakage through D_{min} is negligible; the binary gap is a property of the topology, not a numerical artefact. Body-fitted Navier–Stokes validation on a purpose-designed binary multi-channel manifold is deferred to future work.

References

- [1] Nisong Monyimba, Vincent Pizziconi, and Aurel Coza. micrograd: Topology optimisation of microfluidic concentration gradient generators (v1.0.0), 2026. Zenodo: 10.5281/zenodo.20538833.
- [2] T. Borrvall and J. Petersson. Topology optimization of fluids in stokes flow. *International Journal for Numerical Methods in Fluids*, 41(1):77–107, 2003.
- [3] J. Alexandersen, C. S. Andreasen, N. Aage, and O. Sigmund. Topology optimisation for natural convection problems. *International Journal for Numerical Methods in Fluids*, 76(10):699–721, 2014.
- [4] D. S. Makhija and K. Maute. Level set topology optimization of scalar transport problems. *Structural and Multidisciplinary Optimization*, 51:267–285, 2015.
- [5] S. Kreissl, G. Pingen, and K. Maute. Topology optimization for unsteady flow. *International Journal for Numerical Methods in Engineering*, 87(13):1229–1253, 2011.
- [6] George M. Whitesides. The origins and the future of microfluidics. *Nature*, 442:368–373, 2006.
- [7] J. Na, H. Li, P. Yan, X. Li, and X. Gao. An open-source topology optimization modeling framework for the design of passive micromixer structure. *Chemical Engineering Science*, 260:117820, 2022.
- [8] Jørgen S. Dokken, August Johansson, Mikael Ranta, and Matthew W. Scroggs. Automatic finite element code generation for variational problems in FEniCSx. *ACM Transactions on Mathematical Software*, 49(4):1–26, 2023.
- [9] Boyan S. Lazarov and Ole Sigmund. Filters in topology optimization based on Helmholtz-type differential equations. *International Journal for Numerical Methods in Engineering*, 86(6):765–781, 2011.
- [10] A. N. Brooks and T. J. R. Hughes. Streamline upwind/petrov-galerkin formulations for convection dominated flows with particular emphasis on the incompressible navier-stokes equations. *Computer Methods in Applied Mechanics and Engineering*, 32:199–259, 1982.
- [11] Jean Donea and Antonio Huerta. *Finite Element Methods for Flow Problems*. Wiley, 2003.
- [12] F. Feppon. Multiscale topology optimization of modulated fluid microchannels based on asymptotic homogenization. *Computer Methods in Applied Mechanics and Engineering*, 419:116646, 2024.

- [13] Patrick E. Farrell, David A. Ham, Simon W. Funke, and Marie E. Rognes. Automated derivation of the adjoint of high-level transient finite element programs. *SIAM Journal on Scientific Computing*, 35(4):C369–C393, 2013.
- [14] Sebastian Blauth. cashocs: A computational, adjoint-based shape optimization and optimal control software. *SoftwareX*, 24:101577, 2023.
- [15] F. Wang, B. S. Lazarov, and O. Sigmund. On projection methods, convergence and robust formulations in topology optimization. *Structural and Multidisciplinary Optimization*, 43:767–784, 2011.
- [16] Igor A. Baratta, Joseph P. Dean, Jørgen S. Dokken, Michal Habera, Jack S. Hale, Chris N. Richardson, Marie E. Ringøse, Marie E. Rognes, Matthew W. Scroggs, Nathan Sime, and Garth N. Wells. DOLFINx: The next generation FEniCS problem solving environment. *preprint*, 2023.
- [17] Howard C. Elman, David J. Silvester, and Andrew J. Wathen. *Finite Elements and Fast Iterative Solvers: With Applications in Incompressible Fluid Dynamics*. Oxford University Press, 2014.
- [18] M. P. Bendsøe and O. Sigmund. *Topology Optimization – Theory, Methods, and Applications*. Springer, 2003.
- [19] K. Svanberg. The method of moving asymptotes – a new method for structural optimization. *International Journal for Numerical Methods in Engineering*, 24:359–373, 1987.
- [20] GitHub, Inc. GitHub Actions, 2024. Continuous integration and automation platform. Accessed May 2025.
- [21] Carsten Othmer. A continuous adjoint formulation for the computation of topological and surface sensitivities of ducted flows. *International Journal for Numerical Methods in Fluids*, 58(8):861–877, 2008.
- [22] S. Scott Collis and Matthias Heinkenschloss. Analysis of the streamline upwind/Petrov Galerkin method applied to the solution of optimal control problems. Technical Report TR02–01, Rice University, Department of Computational and Applied Mathematics, 2002.
- [23] Siva K. Nadarajah and Antony Jameson. A comparison of the continuous and discrete adjoint approach to automatic aerodynamic optimization. In *38th Aerospace Sciences Meeting and Exhibit*, 2000.
- [24] Erik Andreassen, Anders Clausen, Mattias Schevenels, Boyan S. Lazarov, and Ole Sigmund. Efficient topology optimization in MATLAB using 88 lines of code. *Structural and Multidisciplinary Optimization*, 43(1):1–16, 2011.
- [25] M. B. Giles and N. A. Pierce. An introduction to the adjoint approach to design. *Flow, Turbulence and Combustion*, 65:393–415, 2000.

Table 7: Software quality metrics for **brinkgrad** v1.0.0. LOC breakdown separates core algorithmic code (386 lines) from utilities (297) and extended modules (1,220). Tests and examples are listed separately from the package total.

Metric	Value	Notes
<i>LOC breakdown (non-empty, non-comment)</i>		
Core (solver + adjoint + optimiser)	386	4 files: solver, adjoint, optimizer, gradient_optimizer
Utilities (filter/mesh/postprocess)	297	6 files: utilities, mesh, postprocess, etc.
Extended modules	1,220	12 files: validation, scalability, experimental
Package total (brinkgrad/*.py)	1,903	
Examples (examples/*.py)	206	7 user-facing scripts
Tests (tests/*.py)	35	pytest unit tests
<i>Software quality</i>		
Public functions / classes	68	across 9 modules
Unit tests	10	tests/test_core.py + test_import.py
CI integration tests	7	inline in ci.yml
Total automated tests	18	10 unit + 8 integration, every push
Test coverage	52%	utilities.py; FEniCSx modules via CI pipeline
Runtime dependencies	3	numpy, scipy, matplotlib
Runtime (80×20, 1 iter)	~1 s	8-core laptop
Peak memory (20×5, 50 iter)	29.5 MB	tracemalloc
<i>Reproducibility</i>		
Docker image	Yes	dolfinx/dolfinx:v0.7.3 (pinned)
Reproducibility script	1	run_all.sh (52 lines)
CI workflows	1	.github/workflows/ci.yml
Zenodo archival	Yes	DOI: 10.5281/zenodo.20538833
Open-source licence	MIT	GitHub

Table 8: Reproducibility checklist for **brinkgrad** v1.0.0.

Item	Status	Detail
Open source	✓	MIT licence
DOI archive	✓	Zenodo 10.5281/zenodo.20538833
Version pinning	✓	dolfinx/dolfinx:v0.7.3 (Docker)
Container image	✓	Docker; one-command reproduction
Continuous integration	✓	GitHub Actions; runs on every push
Automated tests	✓	18 tests (unit + integration)
Example scripts	✓	7 scripts in examples/
Dependency list	✓	requirements.txt + setup.py
Numerically reproducible	✓	JSON result files; verified across x86-64 machines

Table 9: Five-layer validation summary for **brinkgrad**.

Layer	Test	Result
Code health	18 automated tests	All pass on every CI push
Gradient correctness	Taylor remainder + direction test	Pearson ≥ 0.80 ; slope ≈ 1.0
Pipeline integration	Mini-optimisation (20 iter)	J decreases monotonically
Mesh robustness	Multiple mesh sizes	Expected non-convex behaviour
Benchmark problems	Linear, step, double-peak	RMSE = 0.058, 0.011, 0.007

Table 10: Adjoint gradient verification. FD remainder slope ≈ 1 ; direction test Pearson ≥ 0.80 . Full gradient analysis in Supplementary S10.

Mesh	FD slope	Pearson corr.	Status
20×5	1.01	0.80	descent direction confirmed
80×20	1.00	0.88	descent direction confirmed

Magnitude analysis in Supplementary S10.

Table 11: Target profile gallery. All runs use the default OC optimiser. Void baseline: single forward solve, $\rho \equiv 1$. RMSE computed on the outlet boundary.

Target	Void RMSE	Optimised RMSE	Notes
Linear ($c^* = y/L_y$)	0.007	0.058	Primary benchmark; dominated by passive diffusion at this Pe
Step ($c^* = \mathbf{1}_{y > L_y/2}$)	0.011	0.011	80×20 , 1,400 iter; passive diffusion dominant
Double-peak (two Gaussians)	0.008	0.007	20×5 , 200 iter; optimiser reduces RMSE below void

Table 12: Anticipated reviewer concerns and brief responses.

Concern	Response
Is this novel?	To the authors’ knowledge, no openly documented FEniCSx implementation combining coupled Brinkman–CD, SUPG-stabilised adjoint, CI verification, and Docker reproducibility has been reported (Table 1).
Why not dolfin-adjoint?	dolfin-adjoint lacks: (i) pre-configured coupled physics, (ii) SUPG at high Pe , (iii) Dockerised reproducibility.
Is the adjoint correct?	Direction test (Pearson ≥ 0.80) and Taylor slope ≈ 1 (Section 4.1); practical convergence in Table 4.
Is it reproducible?	Docker + CI + Zenodo + <code>bash run_all.sh</code> ; numerically reproducible within the pinned environment (Table 8).
What is the runtime?	≈ 1 hour for 1,400 iterations on 80×20 mesh; ≈ 30 min for 600 iterations (RMSE = 0.079).
Can I extend it?	New objective: ~ 8 lines of UFL (Section 3.11). New target: one argument change (Section 3.8).
Why 2D only?	Microfluidic devices are typically 2D; 3D is in the roadmap.

Table 13: Context: outlet RMSE for the linear gradient target on the 80×20 mesh. All values computed with the same forward solver. Heuristic and Stokes-only baselines are not implemented; comparison is against same-code baselines only.

Design	RMSE	Iter	Notes
Void ($\rho \equiv 1$, $V = 1.0$)	≈ 0.007	0	violates $V^* = 0.5$; not a valid baseline
Uniform $\rho = V^*$	0.450	0	Brinkman suppresses flow
Binary-thresholded	0.359	1400	physical upper bound
Short-run ($\beta = 16$)	0.079	600	30 min
Primary ($\beta = 128$)	0.058	1400	main result

Table 14: Optimizer performance across mesh resolutions. The 20×5 mesh is excluded (too coarse; reversed-flow local minima). The 40×10 result shows optimizer sensitivity (plateau after iter 100). The 80×20 primary result uses 1400 iterations (extended β -continuation to 128); the 160×40 validation uses 800 iterations (tuned schedule) (600 iter identical schedule; J monotone, RMSE non-monotone). RMSE measures the global outlet profile error.

Mesh	Elements	Δy (μm)	RMSE	Note
20×5	100	100	—	excluded (too coarse; $N_{\text{out}} = 6$)
40×10	400	50	0.414	optimizer not converged
80×20	1600	25	0.058	Primary result
160×40	6400	12.5	0.091	600 iter, identical; J monotone

Table 15: Binary validation results on the 80×20 mesh. Both binary configurations give identical RMSE, confirming the gap is caused by loss of gray mixing pathways, not by the Brinkman penalisation in fluid regions.

Design	RMSE	Δ RMSE vs gray	Method
Gray Brinkman	0.058	—	Smooth $\alpha(\bar{\rho})$
Binary Brinkman	0.359	+356%	$\bar{\rho} \in \{0, 1\}$, smooth α